

# Hongbin Zhong

📍 Atlanta, US    ✉ hzhong81@gatech.edu    ☎ 470-437-8071    🔗 <https://rjzhb.github.io/>    in Hongbin Zhong  
🌐 rjzhb

## Research Interest

---

Retrieval-Augmented Generation (RAG) Systems  
Data-Centric AI  
Data Systems for Machine Learning  
Data Analytics Systems  
Distributed Machine Learning

## Education

---

**Georgia Institute of Technology** Aug 2024–2029 (expected)  
Ph.D. in Computer Science  
Advisor: Kexin Rong

**Northeastern University** 2020–2024  
B.S. in Computer Science

## Preprints

---

1. **HoneyBee: Efficient Role-based Access Control for Vector Databases via Dynamic Partitioning**  
Hongbin Zhong, Matthew Lentz, Nina Narodytska, Adriana Szekeres, Kexin Rong  
*Submitted to VLDB 2025*
2. **Fast Hypothetical Updates Evaluation**  
*Submitted to top conference demo track*

## Publications

---

1. **FaDE: More Than a Million What-ifs Per Second**  
Haneen Mohammed\*, Alexander Yao\*, Charlie Summers\*, Hongbin Zhong, Gromit Yeuk-Yin Chan, Subrata Mitra, Lampros Flokas, Eugene Wu  
*VLDB 2025*
2. **Accelerating Deletion Interventions on OLAP Workload**  
Haneen Mohammed, Alexander Yao, Lampros Flokas, Hongbin Zhong, Charlie Summers, Eugene Wu  
*ICDE 2024*
3. **PECJ: Stream Window Join on Disorder Data Streams with Proactive Error Compensation**  
Xianzhi Zeng, Shuhao Zhang, Hongbin Zhong, Hao Zhang, Mian Lu, Zhao Zheng, Yuqiang Chen  
*SIGMOD 2024*

## Research Experience

---

**Research Assistant**, Georgia Institute of Technology, Atlanta, GA Aug 2024 – Present  
Advisor: Kexin Rong ; Collaboration: VMware System Group

- Led research on fine-grained access control in **vector databases** for **RAG**, enhancing enterprise data confidentiality.
- Built PostgreSQL/pgvector solutions with row-level security and filtering to optimize storage and retrieval.
- Designed **optimization models** to reduce redundancy and speed up queries through efficient partitioning.

**Research Assistant**, Columbia University, New York City, NY Jul 2023 – Nov 2023  
Advisor: Eugene Wu

- FADE Project - Developed optimization techniques for sparse matrix evaluations, improving performance.
- Applied SIMD and multithreading for sparse data evaluations, reducing disk I/O significantly.

**Research Assistant**, Rutgers University, New Jersey

*June 2023 – Sep 2023*

*Advisor: Dong Deng*

- Implemented baseline methods for data similarity tasks and assisted with running experiments.
- Optimized parallelization for group function tasks in data processing.

**Research Assistant**, Nanyang Technological University / 4Paradigm, Singapore

*Jan 2023 – Jul 2023*

*Advisors: Mian Lu, Shuhao Zhang*

- Developed high-accuracy, low-latency stream processing system for out-of-order data.
- Implemented Bayesian variational inference with transformers for complex data streams.

## Industry Experience

---

**Database Internals Engineer Intern**, InfiniFlow(**vector database startup**)

*Mar 2024 – Apr 2024*

- Improved the mechanism for recording the oldest visible timestamp to avoid unnecessary access to ‘txn\_map’.
- Optimized the cleanup process for bulk deletion of files and records, significantly reducing file I/O operations.

**Full Stack Software Engineer Intern(part-time)**, 4Paradigm

*Feb 2024 – Apr 2024*

- Enhanced AI assistant server performance by refining cache systems, reducing system overhead, and improving user access speed.
- Developed backend logic for community features, and implemented timed tasks for data updates using asynchronous programming.

**Backend Software Engineer Intern**, Meituan, Beijing

*Apr 2022 – Sep 2022*

- Contributed to the Meituan App’s short video project by building foundational features.
- Developed a data reporting pipeline using Kafka and Hive to support recommendation algorithms.
- Improved user experience under poor network conditions by implementing periodic data refreshes through scheduled tasks.

## Technologies

---

**Languages:** C++, C, Java, Python, C#, SQL

**Technologies:** CUDA, Compiler, Database, Deep Learning System, .NET, OS